
Database Systems and Administration
CS3543

Decision-Support Systems: Data Warehousing and advanced SQL

Suprio Ray

University of New Brunswick, Fredericton, Canada

Acknowledgement

Thanks to Ramakrishnan and Gehrke, C. J. Date, Silberchatz, Korth and Sudarshan and many articles for some of the materials in these slides.

Introduction

- Increasingly, organizations are analyzing current and historical data to identify **useful patterns** and **support business strategies**.
- Emphasis is on **complex, interactive, exploratory analysis** of very large datasets created by integrating data from across all parts of an enterprise; data is fairly static.
 - Contrast such **On-Line Analytic Processing (OLAP)** with traditional **On-line Transaction Processing (OLTP)**: mostly long queries, instead of short update transactions.

OLTP vs. OLAP

	OLTP	OLAP
users	clerk, IT professional	knowledge worker, executives
function	day to day operations	decision support
DB design	application-oriented	subject-oriented
data	current, up-to-date detailed, flat relational isolated	historical, summarized, multidimensional integrated, consolidated
usage	repetitive	ad-hoc
access	read/write index/hash on prim. key	lots of scans
unit of work	short, simple transaction	complex query
# records accessed	tens	millions
#users	thousands	hundreds
DB size	100MB-GB	100GB-TB
metric	transaction throughput	query throughput, response

Decision Support Systems

- Decision-Support systems are used to make **business decisions** often based on data collected by **on-line transaction-processing** systems.
- Examples of business decisions:
 - What items to stock?
 - What insurance premium to change?
 - Who to send advertisements to?
- Examples of data used for making decisions
 - Retail sales transaction details
 - Inventory transaction details
 - Customer profiles (income, age, sex, etc.)




Decision-Support Systems: Overview

- **Data Warehousing:** Consolidate data from many sources in one large repository.
 - Loading, periodic synchronization of replicas (ETL)
 - Semantic integration
- **OLAP:** View data “dimensionally”, explore, navigate
 - Complex SQL queries and views.
 - Queries based on spreadsheet-style operations and “multidimensional” view of data.
 - Interactive and “online” queries.
- **Data Mining:** Exploratory search and analysis for interesting trends and anomalies.


Data Warehousing

- Integrated data spanning long time periods, often augmented with summary information.
- Several gigabytes to terabytes common.
- **Need:** interactive response times expected for complex queries; ad-hoc updates uncommon.

EXTERNAL DATA SOURCES



EXTRACT TRANSFORM LOAD (ETL) REFRESH



Metadata
Repository




DATA
WAREHOUSE

SUPPORTS


DATA
MINING



OLAP



Data Warehousing (Cont.)



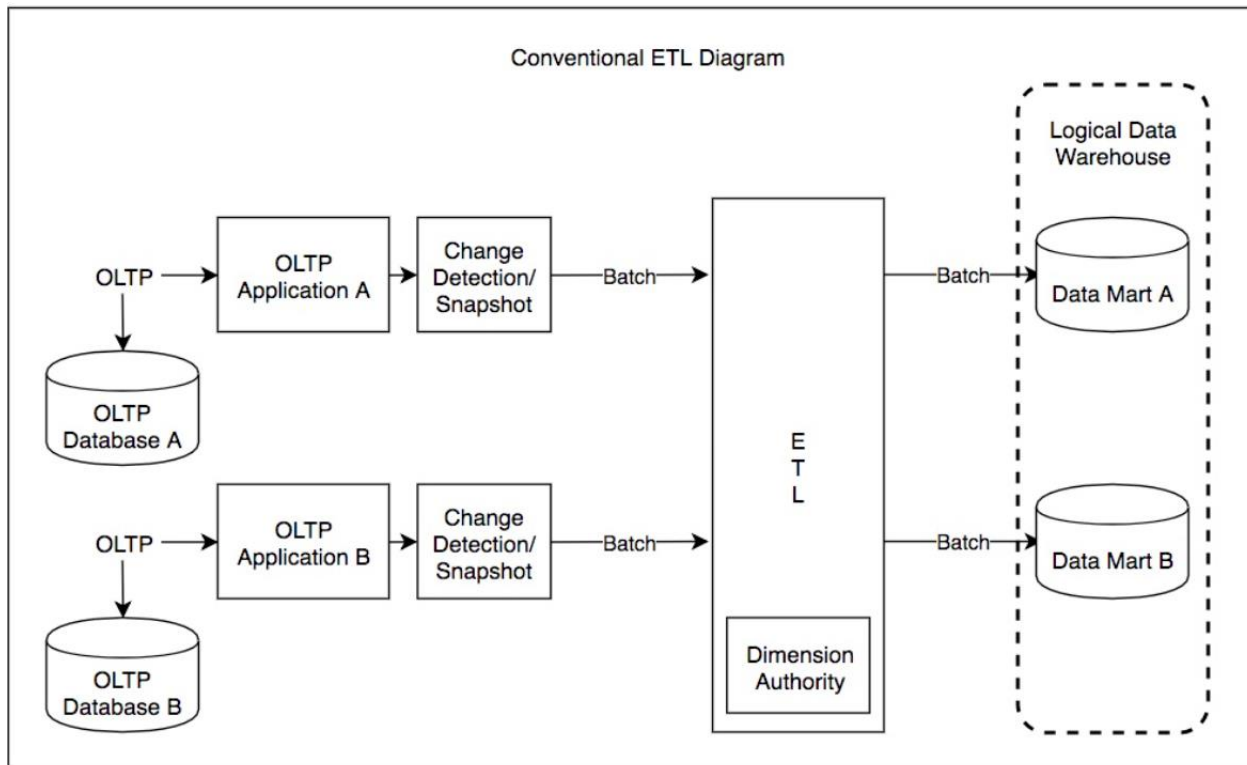
- A data warehouse is a **repository of information** gathered from multiple sources.

Data Warehousing (Cont.)

- Provides a single **consolidated interface** to data
- Data stored for an extended period, providing **access to historical data**
- Running large queries at the warehouse ensures that OLTP systems are not affected by the decision-support workload.
- Data/updates are **periodically downloaded** from online transaction processing (OLTP) systems.
 - Typically, download happens each night.
 - Data may **not be completely up-to-date**, but is recent enough for analysis.

Data Warehousing (Cont.)

- Data/updates are **periodically downloaded** from online transaction processing (OLTP) systems.
 - Typically, download happens each night.
 - Data may **not be completely up-to-date**, but is recent enough for analysis.



Data Warehousing Issues

- **Semantic Integration:** When getting data from multiple sources, must eliminate mismatches, e.g., different currencies, schemas.
- **Heterogeneous Sources:** Must access data from a variety of source formats and repositories.
 - Replication capabilities can be exploited here.
- **Load, Refresh, Purge:** Must load data, periodically refresh it, and purge too-old data.
- **Metadata Management:** Must keep track of source, loading time, and other information for all data in the warehouse.

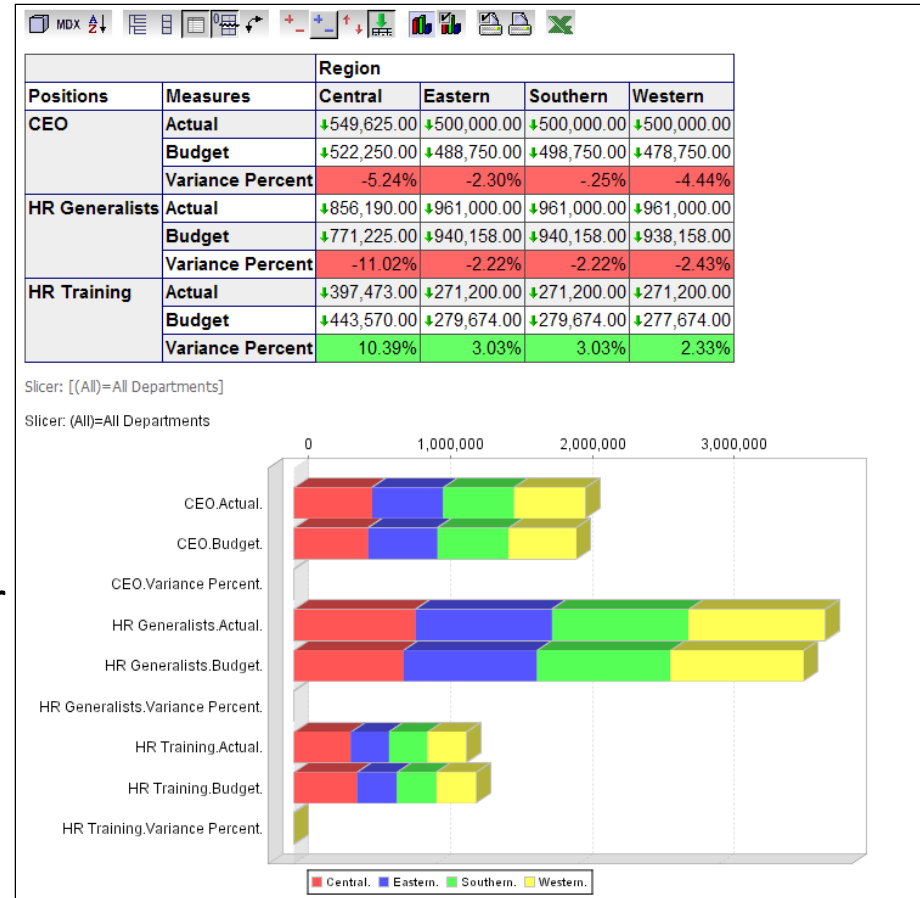
OLAP

- View data “dimensionally”
 - i.e. Sales by region, by channel, by time period

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
13	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35

OLAP

- View data “dimensionally”
 - i.e. Sales by region, by channel, by time period
- Navigate and explore
 - Ad Hoc analysis
 - “Drill-down” from year to quarter
 - Pivot
 - Select specific members for analysis
- Interactive response

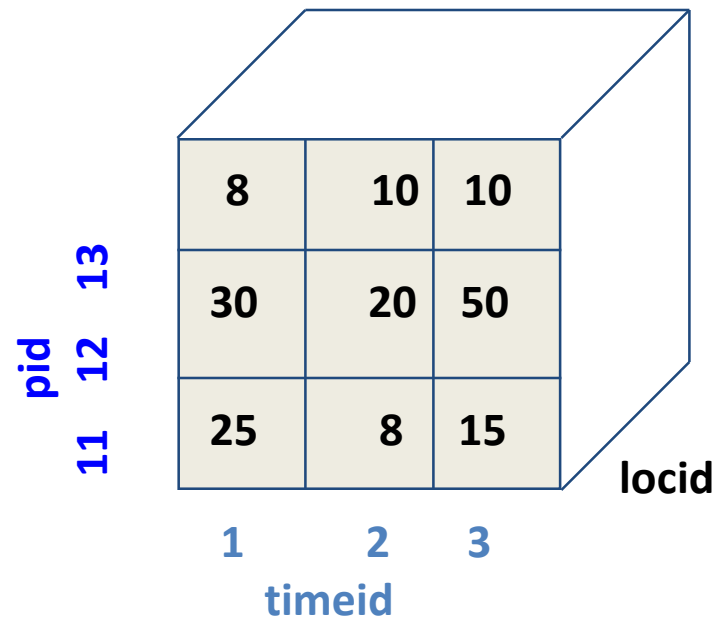


Src: J Hyde

Multidimensional Data Model

- Collection of numeric measures, which depend on a set of dimensions.
 - E.g., measure **Sales**, Revenue, Budget, Expenditure
 - dimensions **Product** (key: pid), **Location** (locid), and **Time** (timeid).

Slice locid=1
is shown:

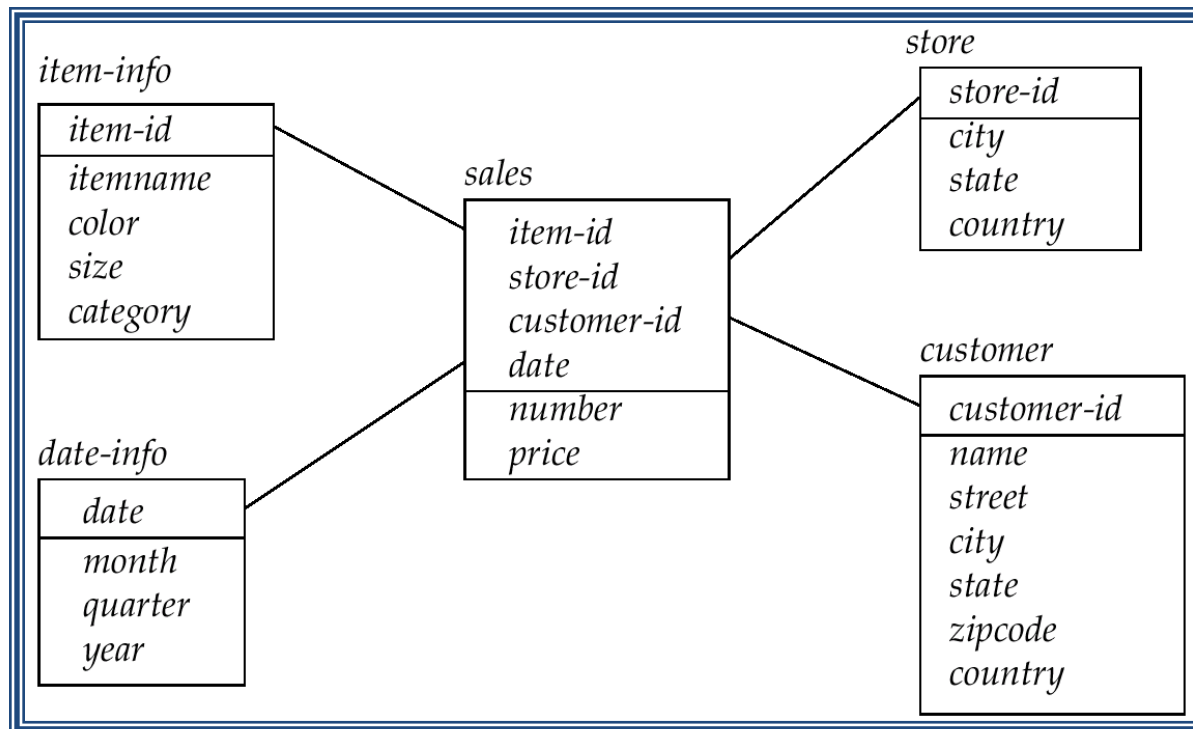


pid	timeid	sales
11	1	8
11	2	10
11	3	10
12	1	30
12	2	20
12	3	50
13	1	25
13	2	8
13	3	15

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
13	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35

OLAP - Star Schema

- The main relation, which **relates dimensions** to **a measure**, is called the **fact table**. Each dimension can have additional attributes and an associated **dimension table**.
 - E.g., **Products(pid, pname, category, price)**
 - Fact tables are *much* larger than dimensional tables.




OLAP: Dimension Hierarchies

- For each dimension, the set of values can be organized in a hierarchy:

PRODUCT



TIME



LOCATION



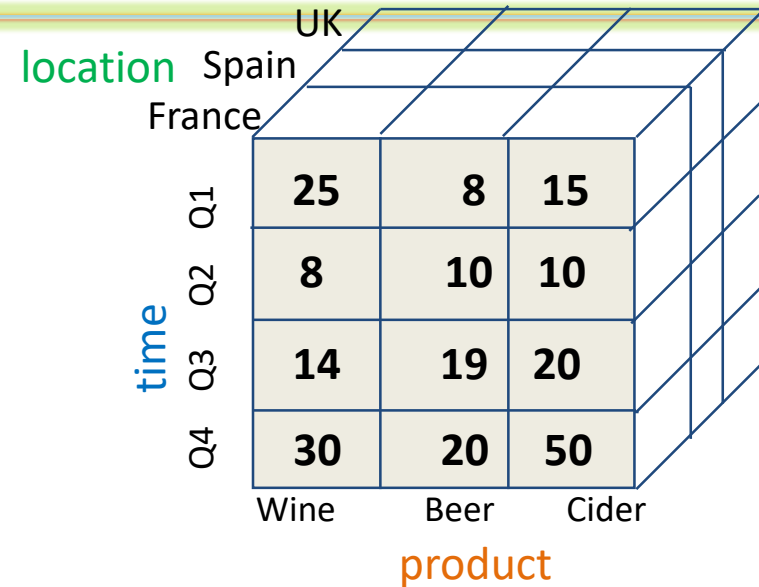
OLAP: Dimension Hierarchies

- Influenced by SQL and by spreadsheets.
- A common operation is to aggregate a measure over one or more dimensions.
 - Find **total** sales.
 - Find **total** sales **for each** city, or for each state.
 - Find **total** sales **for each** quarter, or for each month.
 - Find **top five** products ranked by total sales.

OLAP operations

1. Roll-up: Aggregating at different levels of a dimension hierarchy

- Going from lower to higher in that hierarchy
- E.g., Given total sales **by month**, we can *roll-up* to get total sales **by quarter** (**Month** -> **Quarter**)



A 3D cube diagram representing sales data. The vertical axis is labeled 'time' with values Q1, Q2, Q3, and Q4. The horizontal axis is labeled 'product' with values Wine, Beer, and Cider. The depth axis is labeled 'location' with values France, Spain, and UK. The data values are as follows:

time \ location \ product	Wine	Beer	Cider
Q1	25	8	15
Q2	8	10	10
Q3	14	19	20
Q4	30	20	50



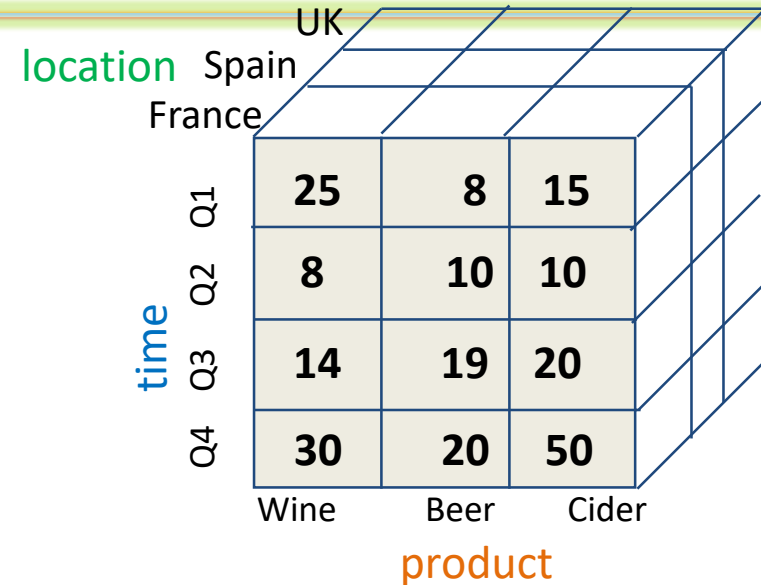
OLAP operations

1. Roll-up: Aggregating at different levels of a dimension hierarchy


- Going from lower to higher in that hierarchy
- E.g., Given total sales **by month**, we can **roll-up** to get total sales **by quarter** (**Month** -> **Quarter**)

2. Drill-down: The inverse of roll-up: Going from higher to lower in the aggregation hierarchy

- E.g., Given total sales **by quarter**, can **drill-down** to get total sales **by Month** (**Quarter** -> **Month**)



		location		
		UK	Spain	France
time	Q1	25	8	15
	Q2	8	10	10
	Q3	14	19	20
	Q4	30	20	50
		Wine	Beer	Cider
		product		



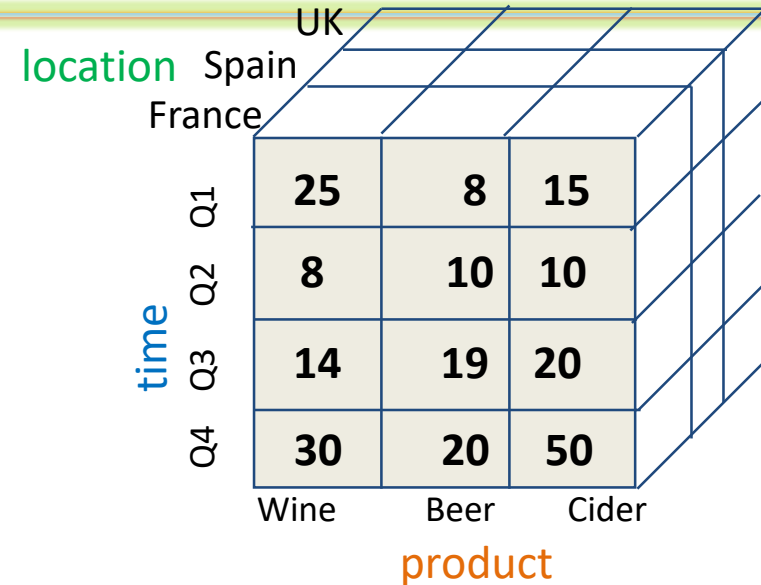
OLAP operations

1. Roll-up: Aggregating at different levels of a dimension hierarchy

- Going from lower to higher in that hierarchy
- E.g., Given total sales **by month**, we can **roll-up** to get total sales **by quarter** (**Month** -> **Quarter**)

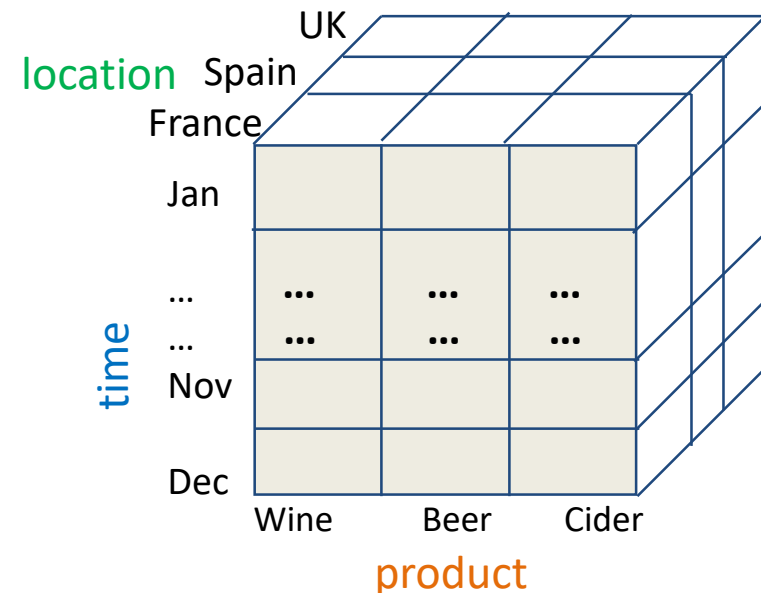
2. Drill-down: The inverse of roll-up: Going from higher to lower in the aggregation hierarchy

- E.g., Given **total sales by quarter**, can **drill-down** to get **total sales by Month** (**Quarter** -> **Month**)



A 3D cube diagram illustrating the roll-up operation. The vertical axis is labeled 'time' with values Q1, Q2, Q3, and Q4. The horizontal axis is labeled 'product' with values Wine, Beer, and Cider. The depth axis is labeled 'location' with values France, Spain, and UK. The data is aggregated by quarter.

time	Wine	Beer	Cider
Q1	25	8	15
Q2	8	10	10
Q3	14	19	20
Q4	30	20	50



A 3D cube diagram illustrating the drill-down operation. The vertical axis is labeled 'time' with values Jan, ..., Nov, Dec. The horizontal axis is labeled 'product' with values Wine, Beer, and Cider. The depth axis is labeled 'location' with values France, Spain, and UK. The data is aggregated by month.

time	Wine	Beer	Cider
Jan			
...
...
Nov			
Dec			

OLAP operations

3. Slicing: Equality selections on one or more dimensions.

- Selecting a single dimension results in a new sub-cube creation
- E.g., total sales by **location=France**.

location
= France

time	Q1	25	8	15
	Q2	8	10	10
	Q3	14	19	20
	Q4	30	20	50
		Wine	Beer	Cider
		product		

4. Dicing: Range selections on one or more dimensions

- Selecting a sub-cube by selecting two or more dimensions
- E.g., total sales by **time** in (Q1, Q2) and **product** in (Wine, Beer).

location	UK			
	Spain			
	France			
time	Q1	25	8	
	Q2	8	10	
		Wine	Beer	
		product		

OLAP operations

5. Pivoting: Aggregation on selected dimensions. Then rotating the current view to get a new view.

- E.g., Pivoting on Location and Time yields cross-tabulation:

size:

all

color

	dark	pastel	white	Total
skirt	8	35	10	53
dress	20	10	5	35
shirt	14	7	28	49
pant	20	2	5	27
Total	62	54	48	164

item-name

Illustration: flat data table

<i>item-name</i>	<i>color</i>	<i>number</i>
skirt	dark	8
skirt	pastel	35
skirt	white	10
skirt	all	53
dress	dark	20
dress	pastel	10
dress	white	5
dress	all	35
shirt	dark	14
shirt	pastel	7
shirt	white	28
shirt	all	49
pant	dark	20
pant	pastel	2
pant	white	5
pant	all	27
all	dark	62
all	pastel	54
all	white	48
all	all	164


Cross Tabulation of *sales* by *item-name* and *color*

size: all

<i>item-name</i>	<i>color</i>			
	dark	pastel	white	Total
skirt	8	35	10	53
dress	20	10	5	35
shirt	14	7	28	49
pant	20	2	5	27
Total	62	54	48	164

- The table above is an example of a **cross-tabulation** (or **cross-tab**) also referred to as a **pivot-table**. *In general, a cross-table is a table where values for **one attribute form the row headers**, values for **another attribute form the column headers**, and the values in an individual cell are derived as above.*
- A cross tab with summary rows/columns can be represented by introducing a special value all **to represent subtotals**.

Another Sample Data Cube



An example: supplier-and-parts database

S#	P#	QTY
S1	P1	300
S1	P2	200
S2	P1	300
S2	P2	400
S3	P2	200
S4	P2	200

S# -> supplier#

P# -> parts#

QTY -> quantity

(This example is
from C. J. Date's
database textbook)

Queries:

- 1) Get the total shipment (of parts) quantity
- 2) Get total shipment quantities by supplier
- 3) Get total shipment quantities by part
- 4) Get the shipment by supplier and part

Cross Tabulations

	P1	P2	Total
S1	300	200	500
S2	300	400	700
S3	0	200	200
S4	0	200	200
	600	1000	1600

Queries:

- 1) Get the total shipment quantity
- 2) Get total shipment quantities by supplier
- 3) Get total shipment quantities by part

Cross Tabulations benefits

	P1	P2	Total
S1	300	200	500
S2	300	400	700
S3	0	200	200
S4	0	200	200
	600	1000	1600

- Display query results as cross tabulations
 - More readable way
 - Formatted as a simple array
 - Two dimensions (supplier and parts)

Cross Tabulations and SQL queries

1. The total shipment quantity:

```
SELECT SUM(QTY) AS TOTQTY
FROM SP
GROUP BY () ;
```

TOTQTY
1600

	P1	P2	Total
S1	300	200	500
S2	300	400	700
S3	0	200	200
S4	0	200	200
	600	1000	1600

Cross Tabulations and SQL queries

2. The total shipment quantities by suppliers

```
SELECT S#, SUM(QTY) AS TOTQTY
FROM SP
GROUP BY (S#) ;
```

S#	TOTQTY
S1	500
S2	700
S3	200
S4	200

	P1	P2	Total
S1	300	200	500
S2	300	400	700
S3	0	200	200
S4	0	200	200
	600	1000	1600

Cross Tabulations and SQL queries

3. The total shipment quantities by parts

```
SELECT P#,SUM(QTY) AS TOTQTY
FROM SP
GROUP BY (P#) ;
```

P#	TOTQTY
P1	600
P2	1000

	P1	P2	Total
S1	300	200	500
S2	300	400	700
S3	0	200	200
S4	0	200	200
	600	1000	1600

Cross Tabulations and SQL queries

4. the shipment by supplier and part

```
SELECT S#, P#, SUM(QTY) AS TOTQTY
FROM SP
GROUP BY (S#,P#) ,
```

S#	P#	TOTQTY
S1	P1	300
S1	P2	200
S2	P1	300
S2	P2	400
S3	P2	200
S4	P2	200

	P1	P2	Total
S1	300	200	500
S2	300	400	700
S3	0	200	200
S4	0	200	200
	600	1000	1600

Considerations with SQL queries

- Formulation of so many similar but **many distinct queries** is tedious
- Executing the queries is expensive
- Make life easier,
 - more efficient computation
- **Single query** to replace **many distinct queries**
 - GROUPING SETS, ROLLUP, CUBE options
 - Added to SQL standard 1999

GROUPING SETS

- Execute several queries simultaneously with a single one

```
SELECT S#, P#, SUM (QTY) AS TOTQTY
FROM SP
GROUP BY GROUPING SETS ( (S#), (P#) ) ;
```

Single results table

Not a relation !!

null ➔ missing information



S#	P#	TOTQTY
S1	null	500
S2	null	700
S3	null	200
S4	null	200
null	P1	600
null	P2	1000

ROLLUP

```
SELECT S#,P#, SUM ( QTY ) AS TOTQTY
FROM SP
GROUP BY ROLLUP (S#, P#) ;
```



```
GROUP BY GROUPING SETS ( ( S#, P# ), ( S# ) , ( ) )
```

S#	P#	TOTQTY
S1	P1	300
S1	P2	200
S2	P1	300
S2	P2	400
S3	P2	200
S4	P2	200
S1	null	500
S2	null	700
S3	null	200
S4	null	200
null	null	1600

ROLLUP

- The quantities have been “roll up” for each supplier
- Rolled up “along **supplier** dimension”

GROUP BY ROLLUP (A,B,...,Z)

(A,B,...,Z)

(A,B,...)

(A,B)

(A)


()

➡ GROUP BY ROLLUP (A,B) is not symmetric in A and B !

- Notice, hierarchical rollups starts with the **primary group A**, followed by the others in hierarchical order. This hierarchy is implied by **the order in which the fields appear** in the GROUP BY clause: (A,B,...,Z)

CUBE

```
SELECT S#, P#, SUM ( QTY ) AS TOTQTY
FROM SP
GROUP BY CUBE ( S#, P# ) ;
```



```
GROUP BY GROUPING SETS ( (S#, P#), ( S# ), ( P# ), ( ) )
```

S#	P#	TOTQTY
S1	P1	300
S1	P2	200
S2	P1	300
S2	P2	400
S3	P2	200
S4	P2	200
S1	null	500
S2	null	700
S3	null	200
S4	null	200
null	P1	600
null	P1	1000
null	null	1600

CUBE

- Confusing term CUBE (?)
 - Derived from the fact that in multidimensional terminology: data values are stored in cells of a multidimensional array or a hypercube
 - The actual physical storage **may differ**
 - In our example
 - cube has just two dimensions (supplier, part)
- Means “group” by **all possible subsets** of the set {A, B, ..., Z }
 - If there are k dimensions, we have 2^k possible SQL GROUP BY queries that can be generated

The CUBE Operator

- Generalizing the previous example, if there are k dimensions, we have 2^k possible SQL GROUP BY queries that can be generated through pivoting on a subset of dimensions.
- CUBE pid, locid, timeid BY SUM Sales
 - Equivalent to rolling up Sales on all eight subsets of the set {pid, locid, timeid}; each roll-up corresponds to an SQL query of the form:

Lots of work on optimizing the CUBE operator!

```
SELECT SUM(S.sales)
FROM   Sales S
GROUP BY grouping-list
```